



# **NLP BASICS**

## **Introduction to Word Vectors**

**How do we represent the meaning of a word in a machine?**

# Three eras of NLP

## A brief history of NLP

### Symbolic 1940 - 2000

Rule-based systems  
Formal grammars

Lexicon, ontologies,  
grammars

### Statistical Learning 1990 - 2010

Statistical learning theory,  
Graphical probabilistic  
models (e.g. LSA, HMM)

Annotated datasets

### Deep Learning 2010 - now

Deep learning based  
methods,  
Transfer learning in NLP  
(BERT, GPT)

Larger datasets, open  
source libraries  
(hugging face)

# Units in NLP

**Corpus:** A collection of  $m$  documents

$$C = (d_1, d_2, \dots, d_m)$$

(Book(s), wikipedia, all articles of the NYT)

**Document:** A sequence of  $k$  words

$$D = (w_1, w_2, \dots, w_k)$$

(Sentence, paragraph, sequence of paragraphs)

**Token:** A basic unit of a sequence of characters grouped together for processing

(Word, sub-word, character(s))

# Preprocessing

# Preprocessing: Tokenization

- Chunk a character sequence into smaller discrete element(s) (sentence, word, sub-word)

Input: All I know is that I know nothing

Output: ['All', 'I', 'know', 'is', 'that', 'I', 'know', 'nothing']

- The first step in any NLP pipeline
- Challenge: Mostly language-agnostic, but different language systems require other algorithms (e.g. Arabic, Chinese, Korean, Tamil, Urdu, and others)

# Preprocessing: Stemming

- Used to “normalise” word into base form or root form.

Input: celebrates, celebrated, celebrating

Output: celebrate

- Challenge: can produce a root word which may not have any meaning

Input: intelligence, intelligent, intelligently

Output: intelligen

# Preprocessing: Lemmatization

- Reduce inflectional/variant forms to base form

Input: am/are/is

Output: be

Input: car/cars/car's/cars'

Output: car

Input: the boy's cars are different colors

Output: the boy car be different color

- Lemmatization produces the root word which has a meaning.



# Information Extraction

# Information extraction: Part of Speech Tagging

- Annotate each word in a sentence with a part-of-speech marker.

Tag	Description
CC	Coordinating conjunction
CD	Cardinal number
DT	Determiner
EX	Existential there
FW	Foreign word
IN	Preposition or subordinating conjunction
JJ	Adjective
JJR	Adjective, comparative
JJS	Adjective, superlative
LS	List item marker
MD	Modal
NN	Noun, singular or mass
NNS	Noun, plural
NNP	Proper noun, singular
NNPS	Proper noun, plural
PDT	Predeterminer
POS	Possessive ending
PRP	Personal pronoun

Tag	Description
PRP\$	Possessive pronoun
RB	Adverb
RBR	Adverb, comparative
RBS	Adverb, superlative
RP	Particle
SYM	Symbol
TO	to
UH	Interjection
VB	Verb, base form
VBD	Verb, past tense
VBG	Verb, gerund or present participle
VCN	Verb, past participle
VBP	Verb, non3rd person singular present
VBZ	Verb, 3rd person singular present
WDT	Whdeterminer
WP	Whpronoun
WP\$	Possessive whpronoun
WRB	Whadverb

# Information extraction: Part of Speech Tagging

- Annotate each word in a sentence with a part-of-speech marker.

Input: Barack Obama was born in Hawaii and served as the 44th President of the United States.

Output: [Barack]<sub>NNP</sub> [Obama]<sub>NNP</sub> [was]<sub>VBD</sub> [born]<sub>VBN</sub> [in]<sub>IN</sub> [Hawaii]<sub>NNP</sub>  
[and]<sub>CC</sub> [served]<sub>VBD</sub> [as]<sub>IN</sub> [the]<sub>DT</sub> [44th]<sub>JJ</sub> [President]<sub>NN</sub> [of]<sub>IN</sub> [the]<sub>DT</sub>  
[United]<sub>NNP</sub> [States]<sub>NNPS</sub> [.]

- Lowest level of syntactic analysis.

# Information extraction: Named Entity Recognition

- Identify and categorize the text into specific entities
- Entities: who did what, when, where, why

Type	Tag	Sample Categories	Example sentences
People	PER	people, characters	<b>Turing</b> is a giant of computer science.
Organization	ORG	companies, sports teams	The <b>IPCC</b> warned about the cyclone.
Location	LOC	regions, mountains, seas	The <b>Mt. Sanitas</b> loop is in <b>Sunshine Canyon</b> .
Geo-Political Entity	GPE	countries, states, provinces	<b>Palo Alto</b> is raising the fees for parking.
Facility	FAC	bridges, buildings, airports	Consider the <b>Tappan Zee Bridge</b> .
Vehicles	VEH	planes, trains, automobiles	It was a classic <b>Ford Falcon</b> .

**Figure 21.1** A list of generic named entity types with the kinds of entities they refer to.

# Information extraction: Named Entity Recognition

- Identify and categorize the text into specific entities
- Entities: who did what, when, where, why

Input: Barack Obama was born in Hawaii and served as the 44th President of the United States.


Output: [Barack Obama]<sub>PER</sub> was born in [Hawaii]<sub>GPE</sub> and served as the [44th]<sub>ORD</sub> President of the [United States]<sub>GPE</sub>.

# Let's code it!



# NLP modeling framework

- Modeling an NLP task involves estimating the conditional probability

$$p(Y|X)$$


The diagram shows the mathematical expression  $p(Y|X)$ . Below the  $Y$ , a straight arrow points diagonally down and to the left. Below the  $X$ , a curved arrow points diagonally down and to the right. Both arrows converge towards the word "Tokens" which is positioned to the right of the expression.

- Sequence classification** task:  $Y$  is a single label
- Sequence labelling** task:  $Y$  has one label per token
- Sequence prediction** task:  $Y$  is a sequence of tokens
- Structure prediction** task:  $Y$  is a graph or a tree

# Representing text with vectors

- Let's assume "token = word"
- We can represent words in vectors by applying different techniques
  - One-hot encoding
  - Hand-crafted representations
  - Count-based representations
  - Learning-based representations
- Word vectors and word embeddings are used interchangeably



# One-hot encoding


- Every word is represented in an element of a vector
- A word vector consists of 0s everywhere except a single 1 for the word

human	1	0	0	0	...	0
machine	0	1	0	0	...	0
system	0	0	1	0	...	0
for	0	0	0	1	...	0
...	...	...	...	...	...	...
user	0	0	0	0	...	1

- Challenge: Polysemy (river bank, financial institution bank, bank for sitting)

# Hand-crafted representations: WordNet

- WordNet (see also VerbNet, FrameNet)
- Manually annotated lexical database of words and their semantic relationships



`is_similar(dog, lion) = ?`

- Challenge: Requires a lot of human labor, subjectivity of annotators, does not scale

# Word representations using data


- Representing words by their context: **Distributional semantics**

*“You shall know a word by the company it keeps.” (Firth, 1957)*

I deposited my pay check at the **bank** this morning.

I need to visit the **bank** to withdraw some cash for my trip.

I need to transfer some funds at the **bank**.



nearby context words  
will represent  
**bank**


**Idea:** Consider the context of a word to build its vector representation

# Co-occurrence matrix

Text corpus



Distributions of words  
in the text



	human	machine	system	for	...	user
human	0	2	1	0	...	0
machine	2	0	0	0	...	0
system	1	0	0	0	...	1
for	0	0	0	0	...	0
...	...	...	...	...	...	...
user	0	0	1	0	...	0

1. Define the context of a word
2. Count how many times a word co-occurs in this context

# Co-occurrence matrix


The **human** operated the **machine** to complete the task efficiently.

The **user** relied on the **machine's** accuracy to obtain the desired result.

**Humans** and **machine** collaborate to achieve optimal **system** performance.

The **system** was tailored to meet the needs of different **users**.

*word embedding*



	human	machine	system	for	...	user
human	0	2	1	0	...	0
machine	2	0	0	0	...	0
system	1	0	0	0	...	1
for	0	0	0	0	...	0
...	...	...	...	...	...	...
user	0	0	1	0	...	0

## Co-occurrence matrix

Represents the frequency of word co-occurrences within a specified window of text

# Co-occurrence matrix


The **human** operated the **machine** to complete the task efficiently.

The **user** relied on the **machine**'s accuracy to obtain the desired result.

**Humans** and **machine** collaborate to achieve optimal **system** performance.

The **system** was tailored to meet the needs of different **users**.

*word embedding*



	human	machine	system	for	...	user
human	0	2	1	0	...	0
machine	2	0	0	0	...	0
system	1	0	0	0	...	1
for	0	0	0	0	...	0
...	...	...	...	...	...	...
user	0	0	1	0	...	0

## Limitations

Word frequency is very skewed → biases the representations

Good embeddings depends on the size of the corpus

# Solution: Pointwise Mutual Information

Use the probability of co-occurrences instead of absolute counts

$$PMI(w_1, w_2) = \log \frac{p(w_1, w_2)}{p(w_1)p(w_2)}$$

**Idea:** Do words  $w_1, w_2$  co-occur more than if they were independent?

	human	machine	system	for	...	user
human	0	2	1	0	...	0
machine	2	0	0	0	...	0
system	1	0	0	0	...	1
for	0	0	0	0	...	0
...	...	...	...	...	...	...
user	0	0	1	0	...	0



	human	machine	system	for	...	user
human	0	1.54	2.94	0	...	0
machine	1.54	0	0	0	...	0
system	2.94	0	0	0	...	2.54
for	0	0	0	0	...	0
...	...	...	...	...	...	...
user	0	0	2.54	0	...	0

# Solution: Pointwise Mutual Information

Use the probability of co-occurrences instead of absolute counts

$$PMI(w_1, w_2) = \log \frac{p(w_1, w_2)}{p(w_1)p(w_2)}$$

Less sensitive to change in frequent words

**Idea:** Do words  $w_1, w_2$  co-occur more than if they were independent?

	human	machine	system	for	...	user
human	0	2	1	0	...	0
machine	2	0	0	0	...	0
system	1	0	0	0	...	1
for	0	0	0	0	...	0
...	...	...	...	...	...	...
user	0	0	1	0	...	0



	human	machine	system	for	...	user
human	0	1.54	2.94	0	...	0
machine	1.54	0	0	0	...	0
system	2.94	0	0	0	...	2.54
for	0	0	0	0	...	0
...	...	...	...	...	...	...
user	0	0	2.54	0	...	0



## Solution: Pointwise Mutual Information

Use the probability of co-occurrences instead of absolute counts

$$PMI(w_1, w_2) = \log \frac{p(w_1, w_2)}{p(w_1)p(w_2)}$$

Less sensitive to change in frequent words

**Idea:** Do words  $w_1, w_2$  co-occur more than if they were independent?

**Limitations:**

- Very large matrix and word vectors
- Corpus dependency: Words and columns represent words in the corpus!

# Solution: Pointwise Mutual Information

Use the probability of co-occurrences instead of absolute counts

$$PMI(w_1, w_2) = \log \frac{p(w_1, w_2)}{p(w_1)p(w_2)}$$

Less sensitive to change in frequent words

**Idea:** Do words  $w_1, w_2$  co-occur more than if they were independent?

**Limitations:**

- Very large matrix and word vectors
- Corpus dependency: Words and columns represent words in the corpus!

—> A solution: Dimensionality reduction (Singular Value Decomposition)

# Let's code it!



# Learning based representations

**Co-occurrence vectors** are

- Long (#unique words in corpus)
- Sparse (most elements are 0)

**Alternative:** learn vectors which are

- Short (~300 dimensions)
- Dense (most elements are non-zero)

# Learning based representations

- Building a vector for each word
- Learning objective:
  - A word vector is similar to vectors of words that appear in similar contexts
- Similarity measurement: Cosine similarity of two vectors

	d1	d2	d3	d4	d5	d6	d7
<i>dog</i> →	0.6	0.9	0.1	0.4	-0.7	-0.3	-0.2
<i>puppy</i> →	0.5	0.8	-0.1	0.2	-0.6	-0.5	-0.1
<i>cat</i> →	0.7	-0.1	0.4	0.3	-0.4	-0.1	-0.3
<i>houses</i> →	-0.8	-0.4	-0.5	0.1	-0.9	0.3	0.8
<i>man</i> →	0.6	-0.2	0.8	0.9	-0.1	-0.9	-0.7
<i>woman</i> →	0.7	0.3	0.9	-0.7	0.1	-0.5	-0.4
<i>king</i> →	0.5	-0.4	0.7	0.8	0.9	-0.7	-0.6
<i>queen</i> →	0.8	-0.1	0.8	-0.9	0.8	-0.5	-0.9

Word                      Word embedding

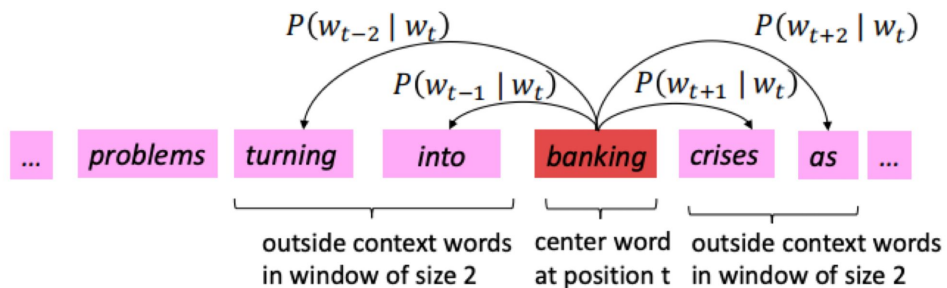
Rozado, PlosOne, 2020

# Learning based representations

## **word2vec: An algorithm for learning word vectors**

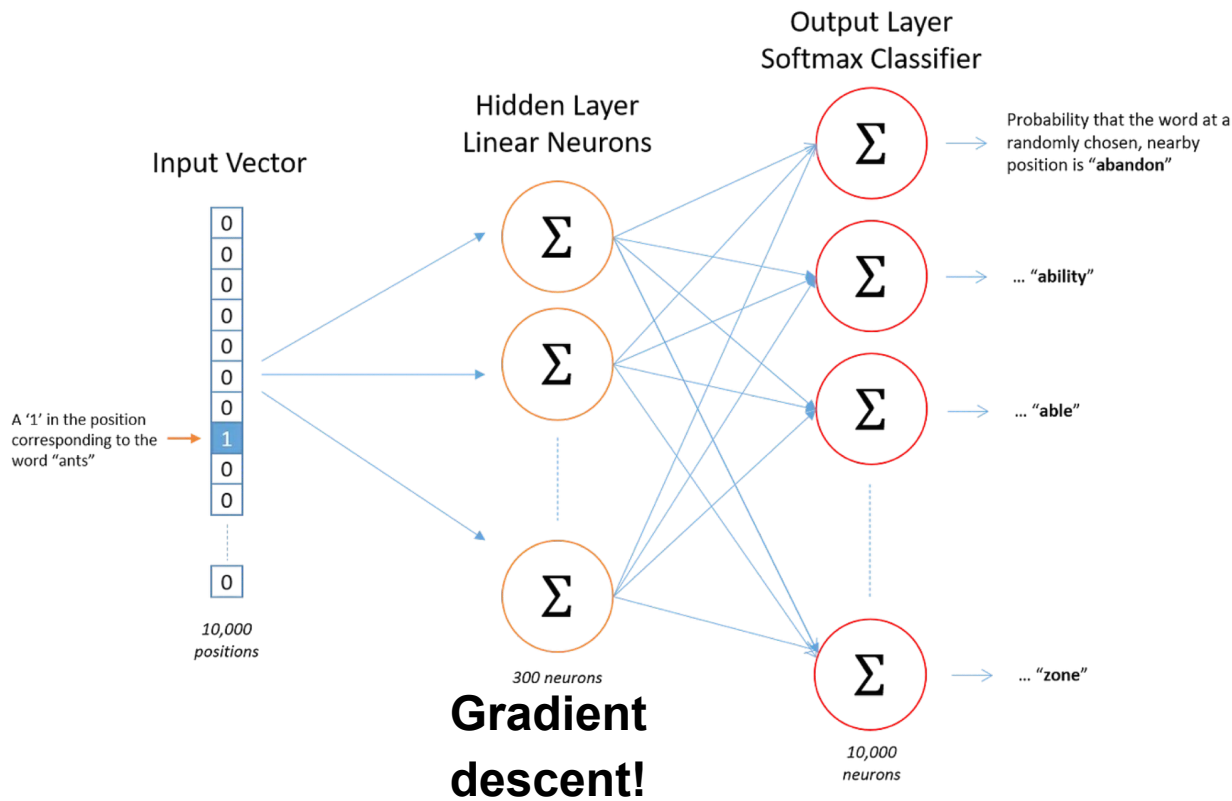
1. Take a large corpus of text
2. Every word is represented by a vector: random initialization
3. Pick a center word  $w$
4. Pick context word  $c$  surrounding the word  $w$
5. Use the similarity of word vectors for  $w$  and  $c$  to compute  $p(c|w)$
6. Repeat to maximize the probability

# word2vec



Hewitt (CS2224N, Stanford)

# The learned weights in the network are the word embeddings





## word2vec: objective function

For each position  $t=1, \dots, T$ , predict context words within a window size  $m$  given a center word  $w_t$  for all parameters  $\theta$ :

Likelihood

$$L_{\theta} = \prod_{t=1}^T \prod_{-m \leq j \leq m} p(w_{t+j} | w_t; \theta)$$

## word2vec: objective function

For each position  $t=1, \dots, T$ , predict context words within a window size  $m$  given a center word  $w_t$  for all parameters  $\theta$ :

Likelihood  $\operatorname{argmax}_{\theta}(L_{\theta}) = \prod_{t=1}^T \prod_{-m \leq j \leq m} p(w_{t+j} | w_t; \theta)$

Maximizing Likelihood

## word2vec: objective function

For each position  $t=1, \dots, T$ , predict context words within a window size  $m$  given a center word  $w_t$  for all parameters  $\theta$ :

Likelihood 
$$\operatorname{argmax}_{\theta}(L_{\theta}) = \prod_{t=1}^T \prod_{-m \leq j \leq m} p(w_{t+j} | w_t; \theta)$$

Objective Function 
$$\operatorname{argmin}_{\theta} \left( -\frac{1}{T} \log(L_{\theta}) \right) = -\frac{1}{T} \sum_{t=1}^T \sum_{-m \leq j \leq m} \log(p(w_{t+j} | w_t; \theta))$$

Maximizing likelihood  $\Leftrightarrow$  Minimizing negative log likelihood

## word2vec: objective function

For each position  $t=1, \dots, T$ , predict context words within a window size  $m$  given a center word  $w_t$  for all parameters  $\theta$ :

Likelihood 
$$\operatorname{argmax}_{\theta}(L_{\theta}) = \prod_{t=1}^T \prod_{-m \leq j \leq m} p(w_{t+j} | w_t; \theta)$$

Objective Function 
$$\operatorname{argmin}_{\theta} \left( -\frac{1}{T} \log(L_{\theta}) \right) = -\frac{1}{T} \sum_{t=1}^T \sum_{-m \leq j \leq m} \log(p(w_{t+j} | w_t; \theta))$$

Maximizing likelihood  $\Leftrightarrow$  Minimizing negative log likelihood

# word2vec: objective function

Calculating the conditional probability

How to calculate the conditional probability?

$$p(w_{t+j} | w_t)$$

**Intuition:** Use two vectors per word

$w$  when the word is a center word

$c$  when the word is a context word

**Familiar?**

$$p(c|w) = \frac{\exp(c^T w)}{\sum_{v \in V} \exp(v^T w)}$$

Dot product measures similarity of  $w$  and  $c$

Normalize over entire vocabulary  $V$

## word2vec: objective function

$$\text{len}(V) = 10000$$

$$v \in V$$

$$\text{argmin}_{\theta}(J_{\theta}) = \text{argmin}_{\theta}\left(-\frac{1}{T}\log(L_{\theta})\right) = -\frac{1}{T} \sum_{t=1}^T \sum_{-m \leq j \leq m} \boxed{\log(p(w_{t+j}|w_t; \theta))}$$

$$\boxed{\log(p(c|w)) = c^T w - \log\left(\sum_{v \in V} \exp(v^T w)\right)}$$

$$p(c|w) = \frac{\exp(c^T w)}{\sum_{v \in V} \exp(v^T w)}$$

Negative log-likelihood is computed for every (at every iteration)

—> Very expensive to compute (for —> 3 million weights)

Solution: Negative sampling

## word2vec: objective function

$$p(c|w) = \frac{\exp(c^T w)}{\sum_{v \in V} \exp(v^T w)}$$

$$\text{softmax}(x_i) = p_i = \frac{\exp(x_i)}{\sum_{j=1}^n \exp(x_j)}$$

- This is a softmax of dot products between context and word vectors
- A softmax transforms a vector of numbers  $x_i$  into a probability distribution  $p_i$
- Frequently used as the activation function in neural network

## word2vec: objective function

- Negative sampling

At each iteration update only a small percentage of the model's weights

Negative log-likelihood is computed over  $K \ll V$  words that are not in the context of  $w$



# Optimization through gradient descent

We have an objective function  $J_{\theta}$  that we want to minimize


Use gradient descent to minimize  $J_{\theta}$

Idea:

For the current value of  $\theta$ , calculate the gradient of  $J_{\theta}$

Take the small step in the direction of the negative gradient

Repeat



Hewitt (CS2224N, Stanford)

---

**Algorithm 1** Skip-Gram Word2vec Training

---

Given a corpus  $C$ , made of a set of unique tokens  $V$ . Hyperparameters: number of negative samples  $K$ , a window size  $l$ , dimension of word vectors  $d$ , learning rate ( $\alpha_t$ )

**Initialize Randomly:**  $\mathbf{W} \in \mathbb{R}^{(V,d)}$  and  $\mathbf{C} \in \mathbb{R}^{(V,d)}$

**for** step  $t$  in  $0..T$  **do**

    ### Step 1: Sampling

    Sample  $s = (w_1, \dots, w_n) \in C$  # a sequence in your corpus (e.g. sentence)

    Sample a pair  $(i, j) \in [1, \dots, n]$  with  $|i - j| \leq l$

    we note  $w = w_i, c = w_j$  represented by vectors  $\mathbf{w}$  in  $\mathbf{W}$  and  $\mathbf{c}$  in  $\mathbf{C}$

    Sample  $N_K = \{v_1, \dots, v_K\} \subset V$  represented by  $\{\mathbf{v}_1, \dots, \mathbf{v}_K\}$  in  $\mathbf{C}$  # Negative samples

    ### Step 2: Compute loss

$$l(\mathbf{W}, \mathbf{C}) = -\sigma(\mathbf{w}, \mathbf{c}) - \frac{1}{K} \sum_{v \in N_K} \log \sigma(-\mathbf{w}, \mathbf{v})$$

    ### Step 3: Parameter update with SGD


$$\mathbf{W}_t = \mathbf{W}_{t-1} - \alpha_t \cdot \nabla l(\mathbf{W}_{t-1}, \mathbf{C}_{t-1})$$

$$\mathbf{C}_t = \mathbf{C}_{t-1} - \alpha_t \cdot \nabla l(\mathbf{W}_{t-1}, \mathbf{C}_{t-1})$$

**end**


# Word2vec: other variations

## Continuous Bag of Words(CBOW) vs. skip-gram



CBOW

Center word prediction




Skip-gram

Context word prediction

# Evaluation of word embeddings

## How can we evaluate the quality of word embeddings?

- Idea: Similar words should have similar vectors
- Visualize word embeddings
  - High dimensional! → PCA or T-SNE
- Similarity measurements
  - Cosine similarity :  $similar(w_1, w_2) = \frac{w_1^T w_2}{\|w_1\| \|w_2\|}$
  - L2 distance :  $similar(w_1, w_2) = \|w_1 - w_2\|$
- Similarity to human judgment (e.g. WordSim353 dataset)




# Word2vec

- word2vec is still widely used
- but “contextualised” language modelling took over in the last years (e.g., BERT)
- Different variations and extensions exist:
  - Continuous Bag of Words (CBOW),
  - GloVe
- Multilingual version by building shared word embeddings across languages
  - FastText
- Limitations:
  - Trained on fixed vocabulary
  - Each token has a unique representation (e.g. the word “bank”)

# Let's code it!



# Tomorrow: Language modelling



# Bibliography and Acknowledgement

Benoit and Sagot 2022, “Algorithms for speech and natural language processing”, MVA course material

Muller, “The Basics of Natural Language Processing” course material

Manning, “Natural Language Processing with Deep Learning” course material

Chris McCormick’s blog: <http://mccormickml.com/2016/04/19/word2vec-tutorial-the-skip-gram-model/>

Mikolov et al., Distributed Representations of Words and Phrases and their Compositionality, 2013, <https://arxiv.org/pdf/1310.4546.pdf>

Mikolov et al. Efficient estimation of word representations in vector space. ICLR Workshop, 2013, <https://arxiv.org/pdf/1301.3781.pdf>

Aguirre et al., A study on similarity and relatedness using distributional and word-net based approaches, In Proceedings of NAACL-HLT, 2009